

Comparative Empirical Evaluation & Algorithmic Guardrail Protocol

Benchmarking G1 (Machaca-Casani et al.) vs. Team Stacked Hybrid Models Across Dual Data-Scale Regimes
Authors: Garv Gulati, Umang Gupta, Nikhil Goyal | Academic ML Research Report

1. Executive Summary & Objective

This report delivers a rigorous, empirical cross-evaluation between the machine learning methodologies established in the **G1 Research Paper (Machaca-Casani et al., Electric Power Systems Research, 2026)** and our team's **Stacked Hybrid Super-Learner architecture**. We address two fundamental research mandates:

- 1. **Evaluating G1 Methods on High-Frequency Telemetry:** Deploying G1's algorithms (Linear Regression, Support Vector Regression with RBF kernel, Random Forest, and Deep MLP 50x50) onto our primary smart home dataset (19,735 records, 50 engineered features).
- 2. **Evaluating Our Model on Small-Sample Urban Data:** Deploying our Stacked Tree Ensemble and Regularized Stacked Hybrid onto G1's 200-sample urban cohort from Piedra Santa (Arequipa, Peru).
- 3. **Algorithmic Degradation Flagging & Automatic Fallback:** Implementing an automated safety guardrail that detects when any model produces a determination coefficient $R^2 < 0$, diagnoses the mathematical failure mechanism, flags the anomaly, and automatically falls back to our robust baseline method.

2. Experiment 1: G1 Methods on the Original High-Frequency Telemetry Dataset (19,735 Records)

The primary smart home dataset comprises 10-minute ZigBee telemetry across 4.5 months, enriched with 50 engineered features (cyclical trigonometric temporal encodings, room-to-outdoor thermal gradients, and occupancy proxies). Evaluated on 5,022 held-out test observations:

Model Architecture	Origin	R ² Score	MAE (Wh)	RMSE (Wh)	MAPE (%)	Status & Diagnostic Flag
Our Team Model: Stacked Hybrid Ensemble (ETR + XGB + HistGBM)	Our Project	0.7504	8.62	11.58	15.53%	OPTIMAL BENCHMARK
Random Forest Regressor (100 Trees)	G1 Method	0.7336	8.93	11.96	16.33%	NORMAL (Passed)
Deep Neural Network (MLP 50x50 Hidden Layers)	G1 Method	0.6400	10.51	13.90	19.07%	NORMAL (Passed)
Support Vector Machine (SVR with RBF Kernel)	G1 Method	0.5331	11.78	15.83	21.09%	NORMAL (Passed)
Linear Regression (Ordinary Least Squares)	G1 Baseline	0.4497	13.18	17.19	23.80%	NORMAL (Underfit)

Key Finding (Large-Scale Regime): In high-volume data regimes (\$N \approx 20,000\$), G1's Deep Neural Network (MLP 50x50) functions stably (\$R^2 = 0.6400\$) because the 11,718 training points are sufficient to constrain its 3,000+ internal synaptic weights. However, our **Stacked Hybrid Ensemble decisively dominates all models (\$R^2 = 0.7504\$)**, reducing MAE by 34.6% relative to Linear Regression.

3. Experiment 2: Our Models vs. G1 Methods on the 200-Sample Urban Cohort Dataset

The Piedra Santa dataset features 200 residential records with 9 socio-demographic and environmental predictors. Data partitioning strictly follows G1's 70/30 protocol ($N_{train} = 140$, $N_{test} = 60$) with Min-Max normalization:

Model Architecture	Origin	R ² Score	MAE (kW)	RMSE (kW)	MAPE (%)	Evaluation Status	Operational Action
Support Vector Machine (SVR with RBF Kernel)	G1 Method	+0.0833	0.5447	0.6598	28.28%	PASSED (Optimal)	Deploy as Primary Model
Linear Regression (Ordinary Least Squares)	G1 Method	+0.0524	0.5282	0.6709	27.67%	PASSED (Robust)	Deploy as Baseline / Edge fallback
XGBoost Regressor (Depth=3, L1/L2 Regularized)	Intermediate	+0.0404	0.5603	0.6751	28.01%	PASSED	Constrained tree option
Our Model: Regularized Stacked Ensemble (SVR+LR+XGB)	Our Project	-0.0462	0.5804	0.7049	28.26%	⚠️ FLAGGED (R ² < 0)	Fallback to SVR / Linear Regression
Random Forest Regressor (100 Trees, Unconstrained)	G1 Method	-0.0532	0.5798	0.7073	29.06%	⚠️ FLAGGED (R ² < 0)	Discard: Overfitting small sample
Our Model: Stacked Tree Ensemble (ETR+XGB+HistGBM)	Our Project	-0.1205	0.5830	0.7295	28.36%	⚠️ FLAGGED (R ² < 0)	Fallback to SVR / Linear Regression
Deep Neural Network (MLP 50x50 Hidden Layers)	G1 Method	-0.7941	0.7685	0.9231	44.24%	⚠️ CRITICAL COLLAPSE	Discard: Severe over-parameterization

4. Root-Cause Analysis: Why Complex Models Fail (R² < 0) on Small Datasets

A negative coefficient of determination ($R^2 < 0$) occurs when the residual sum of squares exceeds the total variance of the target variable:

$$R^2 = 1 - \left[\frac{\sum (y_{true} - y_{pred})^2}{\sum (y_{true} - y_{mean})^2} \right] < 0 \implies \text{MSE}(\text{Model}) > \text{Var}(\text{Target})$$

An $R^2 < 0$ signifies that the model performs **worse than a naive horizontal line predicting the sample mean** ($\bar{y}$). Our empirical diagnosis isolates two distinct failure modes:

- Deep MLP Collapse ($R^2 = -0.7941$):** The dual 50-neuron MLP contains over 3,100 trainable synaptic weights ($10 \times 50 + 50 \times 50 + 50 \times 1$). Optimizing 3,100 parameters on only 140 training observations gives an extreme degree-of-freedom deficit ($N / p \approx 0.045$). The optimizer fits the training data perfectly by memorizing local noise, resulting in catastrophic loss explosion on unseen test points.
- Tree Ensemble Degradation ($R^2 = -0.0532$ to -0.1205):** Unconstrained decision trees recursively bisect the 9-dimensional feature space, isolating small leaf clusters. On small samples with 10.5% behavioral outliers (e.g., midnight AC spikes), tree algorithms create high-variance decision boundaries that over-predict nominal test households.

- **Why SVR and Linear Regression Succeed:** SVR utilizes structural risk minimization bounded by an ϵ -insensitive tube ($C=1.0, \epsilon=0.1$). Errors smaller than ϵ incur zero penalty, and the RBF kernel enforces global smoothness, preventing the decision surface from contorting around localized outliers. Linear Regression is structurally restricted to $p+1 = 10$ parameters, preventing variance runaway.

5. Automated Guardrail Protocol: "Flag and Fallback to Robust Method"

To guarantee production stability, our predictive engine implements an automated architectural switchboard:

OPERATIONAL GUARDRAIL LOGIC:
IF (Sample Count $N < 1,000$) **OR** (Validation $R^2 < 0$):
1. Trigger Alert: [FLAG: MODEL_DEGRADATION_DETECTED]
2. Disable Deep Neural Networks & Complex Tree Stacking.
3. **FALLBACK TO OUR ROBUST METHOD:** Route inference through **Support Vector Regression (SVR with RBF kernel)** or **Linear Regression (OLS)**.
ELSE IF (Sample Count $N \geq 10,000$):
1. Activate **Our Stacked Hybrid Ensemble (ETR + XGB + HistGBM)**.
2. Enable Two-Tier Anomaly Detection (Isolation Forest + Dynamic $3 \times \text{IQR}$ Residuals).

6. Summary Comparison Matrix

Evaluation Dimension	Small-Scale Regime (N = 200, G1 Dataset)	Large-Scale Regime (N = 19,735, Our Dataset)
Recommended Production Model	Support Vector Machine (SVR) / Linear Regression	Our Stacked Hybrid Super-Learner (ETR+XGB+HistGBM)
Achieved R ² Score	+0.0833 (SVR) / +0.0524 (LR)	0.7504 (+0.63 pp above base paper)
Failed / Flagged Architectures	Deep MLP (-0.7941), Random Forest (-0.0532)	None (All achieve positive R ² > 0.44)
Primary Failure Catalyst	Over-parameterization & outlier sensitivity	Underfitting (Linear models fail to capture non-linearities)
Operational Action	Enforce linear/kernel fallback guardrail	Deploy full stacked ensemble with anomaly diagnostics