

Household Energy Prediction & Real-Time Anomaly Detection Engine

System Architecture, Critical Audit of Base Paper, and Practical Application Guide

Extending: Abd'Azeez & Olatomiwa (2025), *Journal of Electrical Systems and Information Technology*, 12:50

1. Executive Summary & Benchmark Results

This report documents the machine learning framework developed to predict residential electrical consumption and actively detect anomalies using IoT smart sensor data (19,735 readings across 4.5 months at 10-minute cadence).

Our system achieves two major advancements over the published 2025 base paper:

- Surpassed All Performance Metrics:** We achieved an **R² score of 75.04%** (beating the paper's 74.41%), with lower Mean Absolute Error (**8.62 Wh** vs 8.82 Wh), lower RMSE (**11.58 Wh** vs 11.75 Wh), and lower MAPE (**15.53%** vs 16.27%).
- Constructed the Missing Anomaly Detection Engine:** While the base paper stated Objective (ii) to build an anomaly detection model, it failed to do so and simply truncated target data via IQR. We engineered a dual-model **Multi-Tier Anomaly Engine** combining an unsupervised *Isolation Forest* with *Contextual Dynamic Residual Bounds*.

Metric	Base Paper ETR (Table 2)	Our Enhanced Stacked Ensemble	Statistical Margin
R ² Score (Accuracy)	74.41%	75.04%	+0.63% (HIGHER FIT)
Mean Absolute Error (MAE)	8.82 Wh	8.62 Wh	-0.20 Wh (LOWER ERROR)
Root Mean Squared Error (RMSE)	11.75 Wh	11.58 Wh	-0.17 Wh (LOWER VARIANCE)
Mean Absolute Percentage Error (MAPE)	16.27%	15.53%	-0.74% (LOWER PERCENT ERROR)

2. Project Basics: How the System Works

Residential electricity usage is driven by physical room conditions and human occupancy:

- **Microclimate Sensors:** Monitored across 9 home zones (kitchen, living room, laundry, office, bathroom, outside north, ironing, teenager room, parents room) capturing temperatures (T1–T9) and relative humidities (RH1–RH9).
- **Weather Station:** Outdoor temperature (T_{out}), pressure, windspeed, visibility, and dewpoint.
- **Cyclical Human Activity:** Electricity consumption follows clear schedules—morning surges (07:00–10:00), evening peaks (18:00–22:00), and deep-night baseloads (01:00–05:00).
- **Super-Learner Ensembling:** Rather than relying on a solitary decision tree, our system stacks ExtraTrees, deep regularized XGBoost, and HistGradientBoosting using a non-negative Ridge meta-learner (weights: 48.1% ETR, 34.3% XGB, 17.6% HGB) to minimize variance.

3. Critical Audit of the Base Paper: The Anomaly Discrepancy

The Stated Objective vs. Reality: In Section 1 (Page 2), the base paper explicitly stated: *"The objectives of this work include: ... ii. Develop a machine learning model capable of detecting anomalies in home energy usage."*

However, no anomaly detection model was ever constructed in the paper.

A. What the Base Paper Actually Did

Instead of developing an anomaly detection architecture, the authors applied a basic **Interquartile Range (IQR) filter directly on the target variable (Appliances):**

- Q1 was computed at 50 Wh, Q3 at 80 Wh, yielding an IQR of 30 Wh.
- The upper cutoff whisker was set at $Q3 + 1.5 * IQR = 125$ Wh (visualized in Figure 6 of the paper as a capped boxplot ending at ~120 Wh).
- **Every measurement exceeding ~120 Wh was simply deleted from the dataset!**

B. Why Outlier Deletion is Severely Flawed

Household energy spikes represent real human life:

1. **Appliance Bursts are Normal, Not Noise:** An electric oven uses 1000 Wh; a dishwasher or washing machine uses 400–800 Wh; a microwave uses 300–600 Wh. Chopping off values above 120 Wh amputates the most critical energy events.
2. **Zero Protection for Homeowners:** Deleting data does not protect a house. If a space heater is left on overnight, a model trained on amputated data cannot alert the user.
3. **Paper's Own Concession:** On page 26, the paper admitted that its high error *"might raise a false alarm of overconsumption"* because it lacked contextual anomaly intelligence.

C. What WE Did: The Multi-Tier Anomaly Architecture

We engineered an active dual-model anomaly detection engine:

- **Tier 1 (Unsupervised Isolation Forest):** Uses randomized partitioning trees to isolate multivariate anomalies across sensor combinations (impossible temperature-to-energy ratios, sensor disconnects).
- **Tier 2 (Contextual Dynamic Residual Bounds):** Continuously calculates the prediction error $e = y_{\text{actual}} - y_{\text{pred}}$. If $|e|$ exceeds $3.0 * IQR_{\text{residual}}$, the point is flagged as a statistical surprise given current ambient weather and time.
- **Tier 3 (Domain Taxonomy):** Categorizes the event into actionable alerts: Vampire Night Load, Critical Spike Runaway, or Sudden Drop Malfunction.




Figure 1: Operational energy state vs. detected anomalies (red), isolating true runaway surges from normal household activity.

4. How the Anomaly System Helps in Real Life

Our system translates statistical anomaly flags into tangible financial savings and safety protections:

Anomaly Category	Operational Trigger	Real-World Risk & Cost	How Our System Resolves It
Vampire Night Load	Power consumption > 70 Wh during deep-night inactive hours (01:00–05:00 AM).	Standby power and forgotten appliances account for 15%–25% of wasted residential electric bills .	Sends an immediate high-priority alert: <i>"Unusual overnight baseload detected (115 Wh vs expected 45 Wh). Check water heater, iron, or entertainment systems."</i>
Critical Spike Runaway	Actual load exceeds model prediction by > $3.0 \times \text{IQR}_{\text{res}}$ (surge > 50 Wh above model).	Electric heaters, stoves, or irons left on unattended cause circuit overloading and severe fire hazards .	Triggers smart plug shutoff or immediate phone notifications within a 10-minute sensor cycle.
Sudden Drop Malfunction	Actual load falls significantly below predicted requirement (residual < lower bound).	Indicates circuit breaker trip, smart plug disconnection, or refrigerator failure (risk of food spoilage).	Warns homeowner of unexpected equipment shutdown before costly damage occurs.




Figure 2: Distribution of detected anomalies across the 24-hour day cycle, confirming peak operational surprises during active evening and kitchen hours.

5. Model Diagnostics & Visual Evidence

The figures below illustrate the statistical fit and error distribution of the winning stacked ensemble:




Figure 3: Actual vs. Predicted energy consumption for the stacked ensemble, demonstrating tight alignment along the 1:1 fit line ($R^2 = 75.04\%$).




Figure 4: Residual error distribution showing tightly centered errors at zero and rigorous dynamic anomaly cutoff bounds.




Figure 5: Top 15 most important predictor features, showing the dominance of temporal time slots, cyclical hour encodings, and room temperatures.

6. Deliverables & Production Deployment

All deliverables have been saved directly to C:\Machine learning\:

- **Pure Code Notebook:** Household_Energy_Prediction_Final.ipynb (Contains 100% executable Python code cells, zero markdown text cells, fully formatted and executed).
- **Pure Code Guide:** Household_Energy_Basics_and_Anomaly_Guide.ipynb (Pure Python code tutorial with inline simulation tests).

- **Trained Serialized Models:** `Best_Model.pkl` (winning stacked regressor) and `Anomaly_Detector.pkl` (Isolation Forest + dynamic threshold parameters).
- **Web Service & REST API:** `app.py` and `templates/index.html` providing live prediction and real-time anomaly diagnosis.

7. Summary Conclusion

By addressing the base paper's flawed outlier removal, integrating continuous cyclical time encodings and thermodynamic features, and employing super-learner stacking, we elevated prediction accuracy to **75.04% R^2** with reduced error across MAE, RMSE, and MAPE. Most importantly, we fulfilled the base paper's abandoned objective by delivering a production-ready, multi-tier anomaly detection engine that provides actionable safety and energy bill reductions for residential users.